

# Construction and Analysis of LeNet 5 Model From Strach on Mnist Dataset

ADHIP BHATTARAI<sup>1</sup> and BISHAL RIJAL<sup>2</sup>

<sup>1</sup>IOE Thapathali Campus, Kathmandu (e-mail: adhipbh200@gmail.com)

<sup>2</sup>IOE Thapathali Campus, Kathmandu (e-mail: bishalrijal5467@gmail.com)

Corresponding author: Adhip Bhattarai (e-mail: adhipbh200@gmail.com) and Bishal Rijal (e-mail: bishalrijal5467@gmail.com).

**ABSTRACT** In the realm of deep learning, convolutional neural networks (CNNs) have revolutionized the field of image recognition and classification. LeNet-5, a pioneering architecture designed by Yann LeCun et al., laid the foundation for modern CNNs and remains a fundamental model for understanding their structure and functionality. This article presents an in-depth exploration of the LeNet-5 architecture and demonstrates its implementation on the MNIST dataset—a widely used benchmark in the realm of handwritten digit recognition. Through this tutorial, readers will gain insights into the layers, concepts, and principles that constitute LeNet-5, fostering a deeper understanding of how CNNs work.

**INDEX TERMS** CNN, deep learning, Lenet-5, MNIST-dataset

## I. INTRODUCTION

CONVOLUTIONAL Neural Networks (CNNs) have emerged as a pivotal advancement in the field of machine learning, particularly in tasks involving image analysis and recognition. The inception of the LeNet-5 architecture by Yann LeCun et al. marked a significant milestone in the evolution of CNNs. LeNet-5 was originally designed for handwritten digit recognition and paved the way for subsequent developments in deep learning. Its elegant yet effective design showcased the power of convolutional layers, pooling operations, and fully connected layers in extracting meaningful features from images and making accurate predictions.

In this article, we embark on a journey to demystify the LeNet-5 architecture, unraveling its intricate layers and mechanisms. Through a step-by-step approach, we'll guide you in constructing the LeNet-5 model from scratch using TensorFlow/Keras and training it on the MNIST dataset. By delving into the architectural nuances, activation functions, and layer interactions, we aim to equip you with a solid grasp of how LeNet-5 functions at its core.

This tutorial assumes a basic understanding of neural networks and machine learning concepts. Whether you're a newcomer seeking to comprehend the inner workings of CNNs or an experienced practitioner aiming to revisit the roots of deep learning, this article will serve as a comprehensive resource to enhance your knowledge and practical skills.

In the following sections, we'll dive into the architectural components of LeNet-5, explain its layer-by-layer construc-

tion, demonstrate the implementation process, and finally, present the results of applying LeNet-5 to the MNIST dataset.

## II. METHODOLOGY

### A. THEORY


Convolutional neural networks (CNNs) and related architectures are the go-to architecture for computer vision tasks. They are heavily inspired by our own visual cortex. Computer scientists once studied the human brain and then transferred that knowledge to computers. Around the 1960s, two researchers, Hubel and Wiesel, performed a series of experiments on the brains of half-awake cats. They showed different shapes to the cats and noticed that one particular neuron, now known as the "Hubel and Wiesel neuron," was active for specific cells. During their research, they also discovered two types of cells in the visual cortex: simple cells, which detect features, and complex cells, which have a larger receptive field and can summarize the output of simple cells. This groundbreaking work laid the foundation for the development of Convolutional Neural Networks, computer scientist took inspiration from this work revolutionizing computer vision and pattern recognition tasks.

The main useful features of CNNs are shift invariance and parameter sharing which make them useful for tasks like image classification, object detection, segmentation and many more.


Typical CNN architecture consists of a stack of convolution, pooling, and activation layers, and finally fully

connected dense layers. Different permutations and combinations of these layers, kernel sizes, pooling operations, regularization, normalization, and a few other techniques have resulted in different architectures.


The network has 5 layers with learnable parameters and hence named LeNet-5. [1] It has three set of convolution layers with a combination of average pooling. After the convolution and average pooling layers, we have two fully connected layers. At last, a softmax classifier which classifies the images into respective class. The input of this model



is a 32\*32 grayscale image hence the number of channels is one. Then, the first convolution operation with the filter size 5\*5 and it contains 6 filters. As a result, output get a feature map of size 28\*28\*6. Here, the number of channels is equal to the number of filters applied. After the first pooling



operation, we apply the average pooling and the size of the feature map is reduced by half. Then, next convolution layer with sixteen filters of size 5\*5 also the feature map changed output to 10\*10\*16. The output size is calculated in a similar manner then average pooling or subsampling layer is applied which again reduce the size of the feature map by half i.e 5\*5\*16. Then, it has a final convolutional layer of size 5\*5



with 120 filters as shown in figure. Leaving the feature map size 1\*1\*120. After which flatten result is 120 values. It also has a fully connected layer with eighty-four neurons and at

last, it has an output layer with ten neurons using softmax activation function since the data have ten classes.




FIGURE 1. Architecture of LeNet-5

## B. ARCHITECTURE DETAILS

Layer	# filters / neurons	Filter size	Stride	Size of feature map	Activation function
Input	-	-	-	32 X 32 X 1	
Conv 1	6	5 * 5	1	28 X 28 X 6	tanh
Avg. pooling 1		2 * 2	2	14 X 14 X 6	
Conv 2	16	5 * 5	1	10 X 10 X 16	tanh
Avg. pooling 2		2 * 2	2	5 X 5 X 16	
Conv 3	120	5 * 5	1	120	tanh
Fully Connected 1	-	-	-	84	tanh
Fully Connected 2	-	-	-	10	Softmax

## C. CONVOLUTION NEURAL NETWORK (CNN)

Convolutional Neural Network (ConvNet or CNN) is a type of deep learning algorithm that is mainly used in image and video recognition tasks. It works by applying a set of filters to the input data, creating feature maps which are then processed through multiple layers to produce a prediction. The key features of CNNs include their ability to learn hierarchical representations of data, local connectivity, shared weights, and pooling operations. These properties allow CNNs to effectively extract features from images and video frames, making them well-suited for tasks such as object detection, image classification, and segmentation. The role of all these topics in this project is discussed below:

### • Convolution Layer:

The primary purpose of Convolution in the case of a ConvNet, was to extract features from the input image and create a compact representation that was usable for further processing by the rest of the network. Convolution preserves the spatial relationship between pixels by learning image features using small squares of input data. It was a process where a small matrix of numbers (called kernel or filter) was taken and was passed over the image and transformed it based on the values from the filter. In a typical CNN architecture, multiple convolutional layers are stacked together, allowing the network to learn increasingly complex features. The dimensions of the tensors were visualized using the following equation:

$$z^i = w^i \cdot A^{i-1} + b^i \quad (1)$$

$$A^i = g^i \cdot z^i \quad (2)$$

Where,

$z^i = \text{output of the neurons located in layer 1}$

$w^i = \text{weight of neurons in layer 1}$

$g^i = \text{activation function}$

$b = \text{bias in layer 1}$

- **Activation Functions:**

Activation function is a non-linear mathematical function applied to the output of each neuron in the network. The activation function was used to introduce non-linearity into the Network, allowing it to model complex relationships between the inputs and outputs. Neuron could not learn with just a linear function attached to it. A non-linear activation function would let it learn as per the difference with respect to error. Thus, the following activation function had been decided to be used.

- **Softmax Activation Function:**

The softmax activation function is a commonly used activation function in machine learning and deep learning for multiclass classification problems. It is used to convert the output of a neural network into a probability distribution over multiple classes.

The softmax function takes a vector of scores as input and returns a vector of probabilities that sum to 1. Each element of the output vector represents the probability that the input belongs to the corresponding class.

The mathematical formula for the softmax function is as follows:

$$\text{softmax}(z)_i = \frac{e^{z_i}}{\sum_i^N e^{z_i}} \quad (3)$$

Where,

$Z$  is the vector of raw outputs from the neural network

$e$  is Euler's number (approximately 2.71828)

- **Average Pooling:**

Average pooling is a pooling operation that selects the average value within a small region of a feature map. Average pooling reduces the spatial dimensions of the feature maps, making the network more computationally efficient. Average pooling is typically used in combination with multiple convolutional layers to form a typical CNN architecture, and is applied repeatedly to extract increasingly complex features from the input data.

- **Fully Connected Layer:**

A fully connected layer is a type of layer that connects every neuron in the previous layer to every neuron in the next layer. It takes the output from the previous layer, which can be the output from one or more convolutional or pooling layers, and applies a linear transformation to the data, followed by a non-linear activation function, such as ReLU. The result is then passed on to the next layer in the network. The purpose of the fully connected layer is to aggregate and interpret the features learned by the previous layers, allowing the network to make predictions based on high-level abstractions of the input data.

### III. WORKING PRINCIPLE

#### A. DATASET PREPARATION

Consider a dataset:

$$X = (x_{ij})_{m \times n}$$

having  $m$  number of instances with  $n$  number of features. For  $x_{ij}$ ,  $i$  represents the data and  $j$  represents the feature. The MNIST dataset is a collection of pictures of handwritten digits [2]. In this case, " $i$ " would represent the total number of pictures we have, and " $j$ " would represent the different aspects we want to know about each picture. In total, each picture has 785 different aspects, where 784 aspects describe the intensity of individual pixel values in a  $28 \times 28$  grid (that's 784 pixels in total), and the last aspect tells us which number the handwritten digit represents, ranging from 0 to 9.

#### B. DATASET PREPROCESSING


To make sure our analysis of the MNIST dataset is accurate and useful, we need to prepare the data properly. This involves a series of steps that help us get the data ready for processing by machine learning algorithms.

- **Label Encoding:** Unlike typical datasets with categorical variables, the MNIST dataset contains pixel values that represent the intensity of each pixel in grayscale images. These values are already numerical and don't require label encoding. However, since our output is multiclass, we performed one hot encoding to the output variable.
- **Normalization:** Normalization is crucial for ensuring that all the pixel values have a consistent scale. In the case of MNIST, pixel values are in the range of 0 to 255 (representing pixel intensity). So, we divided all the input columns by 255 to normalize the pixel values in the range between 0 and 1.
- **Handling Missing Values:** Fortunately, the MNIST dataset doesn't have missing values, as each image is well-defined with its pixel values.

#### C. FORWARD PROPAGATION

Forward propagation in a Convolutional Neural Network (CNN) is the process through which input data, such as images, is passed through the network's layers to produce an output prediction. Each layer in the CNN performs specific operations, such as convolutions, activations, and pooling, that transform the input data into a more abstract and representative feature space, eventually leading to a final prediction.

The Forward Pass starts from the Convolutional Layer where there are two things to note, one is the input image and the other is the filter or kernel. Now, a kernel or filter is also known as feature detector which is a set of weights that holds some of the most important features in the input image. A Convolution operation is performed on this layer with the input image and the kernels, then the resulting output also



known as a feature map is obtained. The mathematical form of this operation is given as:

$$(I*K)(i,j) = \sum_{i=1}^m \sum_{j=1}^n (I_{(i+m-1)(j+n-1)} * rot180(K_{ij})) + b_j \quad (4)$$

The  $I$  is the input(image) and  $K$  is the kernel(same dimensions) having  $i,j$  as indices which represents single pixel in the input image, kernel, and the produced output.  $m,n$  are used to represent the dimensions (height, width) of the filter or kernel. The  $I_{(i+m-1)(j+n-1)}$  represents the specific pixel value of the input image at the position  $(i+m-1,j+n-1)$ . As  $i$  and  $j$  change, the filter is applied through each of the input pixels. There are two sums, the outer summation  $\sum_{i=1}^m$  is used to slide the kernel over the input image which iterates over  $i$  from 0 to  $m-1$  slides vertically over the input image. Meanwhile, the inner summation  $\sum_{j=1}^n$  is used for element-wise multiplication which iterates over  $j$  from 0 to  $n-1$ , representing the horizontal sliding of the kernel over the image. Finally, a bias  $b_i$  is added for providing flexibility for the network.

#### • Forward Pass in Pooling Layer

Feature map is obtained after convolving the input image and kernel, now the next process is to apply pooling which is done on the pooling layer.

In case of **Average Pooling**, a window slides over the input feature map and the average value within each window is selected as the output.

$$(I*K)(i,j) = \begin{bmatrix} 23 & 14 & 25 \\ 27 & 29 & 35 \end{bmatrix}$$

When applying average pooling of shape  $2*2$ , we get,  
 $\text{AvgPool}((I*K)(i,j)) = [\text{avg}(23,14,27,29) \text{ avg}(14,25,29,35)]$   
 $\text{AvgPool}((I*K)(i,j)) = [23 \ 25]$

#### • Forward Pass in Flattening

Here the output obtained from convolution and Pooling is transformed into a one-dimensional vector so that it can be passed to the Dense or Fully connected Layer.

#### • Forward Pass in Fully Connected Layer

A full connection layer is something the same as a Multi-Layer Perceptron (MLP). In a Fully Connected Layer, each neuron is connected to every neuron in the previous layer, forming a fully interconnected net-

work. The term "Fully Connected" is specifically used in CNNs to distinguish it from other layers we have discussed. Unlike convolutional and pooling layers that share weights across inputs, the Fully Connected Layer assigns unique weights to each input value.


The forward pass on each neuron of the Fully Connected Layer is represented as:

$$\text{output} = \sigma(W_{ij} \cdot X_j + bJ) \quad (5)$$

Where,  $W$  = Weights and  $X$  = inputs from the preceding layer,  $\sigma$  = Activation function. This equation is simple and we have learned it when starting the basics of Neural Networks. Here we are computing the weighted sum of inputs, and adding a bias.

#### D. BACK PROPAGATION IN CNN

Backpropagation in a Convolutional Neural Network (CNN) is the process of updating the network's weights and biases based on the calculated gradients of the loss function with respect to these parameters. This process enables the network to learn from the training data by iteratively adjusting its parameters to minimize the prediction error. [3]



#### • Backward Propagation in Fully Connected Layer

The fully connected layer has two parameters - weight matrix and bias matrix. Calculating the change in error with respect to weights  $-\partial E/\partial W$ .

Since the error is not directly dependent on the weight matrix, the concept of chain rule is used to find this value as:

$$\partial E/\partial W = \partial E/\partial O \cdot \partial O/\partial Z_2 \cdot \partial z/\partial W \quad (6)$$

#### – Change in error with respect to output

Suppose the actual values for the data are denoted as  $y'$  and the predicted output is represented as  $O$ . Then the error would be given by this equation:

$$E = (y' - O)^2/2$$

If differentiate the error with respect to the output, it will get the following equation:

$$\partial E/\partial O = -(y' - O)$$

#### – Change in output with respect to $Z_2$ (linear transformation output)

To find the derivative of output  $O$  with respect to  $Z_2$ , we must first define  $O$  in terms of  $Z_2$ . If you look at the computation graph from the forward propagation section above, you would see that the output is simply the sigmoid of  $Z_2$ . Thus,  $\partial O / \partial Z_2$  is effectively the derivative of Sigmoid. Recall the equation for the Sigmoid function:

$$f(x) = 1 / (1 + e^{-x})$$

The derivative of this function comes out to be:

$$f'(x) = (1 + e^{-x})^{-1} [1 - (1 + e^{-x})^{-1}]$$

$$f'(x) = \text{sigmoid}(x)(1 - \text{sigmoid}(x))$$

$$\partial O / \partial Z_2 = (O)(1 - O)$$

#### – Change in $Z_2$ with respect to Weights

The value of  $Z_2$  is the result of the linear transformation process. Here is the equation of  $Z_2$  in terms of weights:

$$Z_2 = W^T \cdot A_1 + b$$

On differentiating  $Z_2$  with respect to  $W$ , we will get the value  $A_1$  itself:

$$\partial Z_2 / \partial W = A_1$$

After finding the individual derivation, the chain rule is applied to find the change in error with respect to weights:

$$\partial E / \partial W = \partial E / \partial O \cdot \partial O / \partial Z_2 \cdot \partial Z_2 / \partial W$$

$$\partial E / \partial W = -(y' - O) \cdot \text{sigmoid}' \cdot A_1$$

The shape of  $\partial E / \partial W$  will be the same as the weight matrix  $W$ . We can update the values in the weight matrix using the following equation:

$$W_{\text{new}} = W_{\text{old}} - lr * \partial E / \partial W$$

- **Backward Propagation in Convolution Layer** For the convolution layer, It had the filter matrix as our parameter. During the forward propagation process, it randomly initialized the filter matrix. Here, It will now update these values using the following equation:

$$\text{new\_parameter} = \text{old\_parameter} - (\text{learning\_rate} * \text{gradient\_of\_parameter})$$

To update the filter matrix, the gradient of the parameter  $-\partial E / \partial f$  need to be find out. The derivative of  $\partial E / \partial f$  can be defined as :

$$\partial E / \partial f = \partial E / \partial O \cdot \partial O / \partial Z_2 \cdot \partial Z_2 / \partial A_1 \cdot \partial A_1 / \partial Z_1 \cdot \partial Z_1 / \partial f$$

The values of the derivatives can be calculated as:

#### – Change in $Z_2$ with respect to $A_1$

To find the value for  $\partial Z_2 / \partial A_1$ , it need to have the equation for  $Z - 2$  in terms of  $A_1$ :

$$Z_2 = W^2 \cdot A_1 + b$$

On differentiating the above equation with respect to  $A_1$ , we get  $W^T$  as the result :

$$\partial Z_2 / \partial A_1 = W^T$$

#### – Change in $A_1$ with respect to $Z_1$

The next value that need to determine is  $\partial A_1 / \partial Z_1$ .

$$A_1 = \text{sigmoid}(Z_1)$$

This is simply the Sigmoid function. The derivative of Sigmoid would be:

$$\partial A_1 / \partial Z_1 = (A_1)(1 - A_1)$$

#### – Change in $Z_1$ with respect to filter $f$

Finally, there is the need of value for  $\partial Z_1 / \partial f$ .

$$Z_1 = X * f$$

Differentiating  $Z$  with respect to  $X$  will simply give  $X$ :

$$\partial Z_1 / \partial f = X$$

Now replacing all the derivatives, overall change in error with respect to the filter:

$$\partial E / \partial f = \partial E / \partial O \cdot \partial O / \partial Z_2 \cdot \partial Z_2 / \partial A_1 \cdot \partial A_1 / \partial Z_1 \cdot \partial Z_1 / \partial f$$

Once the value for  $\partial E / \partial f$  is figure out, this value is used to update the original filter value:

$$f = f - lr * (\partial E / \partial f)$$

## IV. RESULTS

The mnist Dataset is the standard since ages and we are able to use it in our project so that we can analyze the model properly with the predefined model using Pytorch and conquer the outcomes. This is the sample example of mnist dataset as it contains all the handwritten digits as data upto 10 classes (0 to 9) as shown in Figure 2. The pixel intensity values of the sample give the actual status of data as shown in Figure 3 The complete information of the data in dataset is visualized in Figure 4. The Actual Architecture Used can be visualized from Figure 5 as explained in the methodology of the report. After, building the architecture, the dataset is split into the training and testing dataset which are used in the training and validation as well as verification propose.

Since, We have constructed our own Lenet-5 Model from Strach as well as Lenet-5 Model using Pytorch which is very easy to implement as all the modules are already specified. For the Pytorch's Lenet Model, we obtain the accuracy of 98.722 and Loss of 0.04889 in Just 11 Epochs Which is satisfactory. We can also visualize it from Performance Metrics:

We evaluated both models using standard performance metrics, including accuracy and loss, on both the training and testing datasets. We constructed Graph of Accuracy Graph as in Figure 6 and Loss Graph as in Figure 7.

LeNet-5 Implementation from Scratch:

The Dataset is similarly divided into the training and testing dataset as we construct the lenet model similar to the



Pytorch's model also including loss function - Cross-entropy loss which is absent in original model. After Training For 25 Epochs, Results are complementary and satisfactory as we obtain the training accuracy of 90.1266 and test accuracy of 90.27 as shown in Figure 9. We constructed Loss Graph of Manual Model as shown in Figure 8 which in comparison to pytorch model is noisy but if increase the batch size then the noise is reduced as well as pytorch model contains auto adjust function which will adjust the values so that accuracy and loss is maintained.

The Predicted Output of the model is shown in Figure 10 as we can see that pixel intensity of the model is complementary and the model predicted actual value as right.

The model is performing perfectly with no single incorrect prediction as above. One can run more inference to find the incorrect one, but overall the model's accuracy is acceptable.

It's clearly visible with the attribution scores calculated by Integrated Gradients what part of the input image the model's parameters are giving weightage to as shown in Figure 11

## V. DISCUSSION

The MNIST dataset, a steadfast benchmark in the machine learning landscape, serves as a reliable foundation for evaluating model performance. The dataset's enduring relevance and standardized structure make it an ideal testing ground for assessing the capabilities of digit recognition models. Each image in the dataset encapsulates a handwritten digit from 0 to 9, encompassing ten distinct classes. The pixel intensity values within these images encapsulate the core essence of the handwritten digits, as artfully depicted in Figure 2.

As we delve into model architecture, the venerable LeNet-5 takes the center stage—a neural network design that has withstood the test of time. Figure 5 visually encapsulates this architecture, which we have thoughtfully applied through two distinct lenses: a manual construction from the ground up and the utilization of PyTorch's pre-existing LeNet-5 model. To allow for both model training and validation, we partitioned the dataset into training and testing subsets.

Comparing PyTorch's LeNet-5 with Manual Implementation:

The spotlight shines brightly on PyTorch's LeNet-5, showcasing its exceptional prowess in the realm of digit recognition. With an astounding accuracy of 98.722 achieved in a mere 11 epochs, PyTorch's design, fortified with sophisticated optimization techniques, delivers results that echo its efficiency. Notably, the diminutive loss value of 0.04889 adds further credence to the prowess of PyTorch's machinery. The vivid accuracy and loss graphs (Figures 6 and 7) beautifully illustrate the model's unwavering convergence within a compact temporal window.

On the manual construction front, our LeNet-5 model presents a commendable training accuracy of 90.1266, harmonized with a testing accuracy of 90.27, reached after a patient span of 25 epochs. These achievements, while slightly trailing behind PyTorch's model, demonstrate the potential of a hands-on approach to model creation. It's

important to note that the loss graph (Figure 8) exhibits subtle fluctuations, which might be tamed by refining the batch size. Additionally, the intriguing aspect of PyTorch's self-adjusting capabilities contributes to its stable accuracy and loss trajectories.

## VI. CONCLUSION

In the captivating journey through the corridors of the LeNet-5 architecture applied to the venerable MNIST dataset, a tapestry of insights emerges, weaving together the threads of manual ingenuity and framework sophistication. This fusion of methods engenders a profound understanding of the quintessence of deep learning model development.

The MNIST dataset, standing as a paragon of enduring relevance, enriches our narrative by providing a comprehensive landscape for evaluating model efficacy. Its standardized nature and portrayal of handwritten digits from 0 to 9 establish a fertile ground for our explorations.

The symphony of model architecture resonates through LeNet-5, an emblem of neural network innovation. This musical piece unfolds through two perspectives: the meticulous construction from scratch and the embrace of PyTorch's preconceived masterpiece. The partitioning of the dataset into training and testing domains cultivates an environment where these models can flourish and be scrutinized.

In the spotlight, PyTorch's LeNet-5 shines as a virtuoso. Its swift ascent to an accuracy zenith of 98.722 within a mere 11 epochs serves as a testament to PyTorch's orchestration of optimization techniques. The orchestral ensemble is enriched by a minuscule loss of 0.04889, attesting to the harmony between its components. Figures 6 and 7 illuminate the performance journey with vivid clarity, encapsulating the model's graceful convergence.

In contrast, our manually sculpted LeNet-5 model, though slightly trailing in accuracy at 90.27, underscores the value of crafting a neural masterpiece from scratch. This creation echoes the meticulous craftsmanship of artisans, offering an alternative lens to model development. The animated loss graph in Figure 8 is a visual symphony of its own, illustrating the model's journey through epochs.

This symposium of results and discourse concludes by casting a spotlight on the artistry of deep learning. PyTorch's harmonious orchestration and our artisanal endeavor intersect, offering practitioners a dynamic spectrum of choices. As we tread the path ahead, the harmonization of personalized craftsmanship and the efficiency of frameworks beckons us, promising an orchestra of discoveries waiting to be composed in the realm of machine learning.

## REFERENCES

- [1] Y. LeCun et al., "Lenet-5, convolutional neural networks," URL: <http://yann.lecun.com/exdb/lenet>, vol. 20, no. 5, p. 14, 2015.
- [2] H. Xiao, K. Rasul, and R. Vollgraf, "Fashion-mnist: a novel image dataset for benchmarking machine learning algorithms," arXiv preprint arXiv:1708.07747, 2017.
- [3] C.-C. J. Kuo, "Understanding convolutional neural networks with a mathematical model," Journal of Visual Communication and Image Representation, vol. 41, pp. 406–413, 2016.



**ADHIP BHATTARAI** is a dedicated individual pursuing a Bachelor's degree in Computer Engineering at Tribhuvan University. With a strong passion for machine learning and data science, he is constantly exploring the latest advancements in these fields. Although he may not have notable accomplishments just yet, Adhip's enthusiasm and drive for learning and applying cutting-edge technologies make him a promising and ambitious individual in the world of computer engineering.



**BISHAL RIJAL** is a dedicated individual currently studying Bachelor's in Computer and Technology at Tribhuvan University. Bishal's enthusiasm for research and innovation has led him to undertake various projects and engage in practical applications of his knowledge. He continually seeks to deepen his understanding of the subject matter, staying up-to-date with the latest advancements and trends. With his relentless determination, inquisitive mindset, and expertise in machine learning and data science, Bishal Rijal is poised to make significant contributions to the ever-evolving field of technology.

## VII. APPENDIX

### A. APPENDIX A: FIGURES

://www.overleaf.com/project/64d0e36f16fee799786083e4




FIGURE 2. Original Sample of Mnistnet dataset




FIGURE 3. Pixel Hierarchy of Samples




FIGURE 4. Preprocessed Samples of Mnist Dataset

Layer (type (var_name))	Input Shape	Output Shape	Param #	Trainable
LeNetSVI (LeNetSVI)				True
-Sequential (feature)				True
-Conv2d (0)	[1, 1, 28, 28]	[1, 16, 5, 5]	--	True
-Tanh (1)	[1, 1, 28, 28]	[1, 16, 28, 28]	156	True
-AvgPool2d (2)	[1, 16, 28, 28]	[1, 16, 14, 14]	--	True
-Conv2d (3)	[1, 16, 14, 14]	[1, 16, 10, 10]	2,416	True
-Tanh (4)	[1, 16, 10, 10]	[1, 16, 10, 10]	--	True
-AvgPool2d (5)	[1, 16, 10, 10]	[1, 16, 5, 5]	--	True
-Sequential (classifier)				True
-Lflatten (6)	[1, 16, 5, 5]	[1, 16]	--	True
-Linear (1)	[1, 16]	[1, 120]	48,120	True
-Tanh (2)	[1, 120]	[1, 120]	--	True
-Linear (3)	[1, 120]	[1, 84]	10,164	True
-Tanh (4)	[1, 84]	[1, 84]	--	True
-Linear (5)	[1, 84]	[1, 10]	658	True
Total params: 61,786				
Trainable params: 61,786				
Non-trainable params: 0				
Total multi-adds (MB): 0.42				
Input size (MB): 0.00				
Forward/Backward pass size (MB): 0.05				
Params size (MB): 0.25				
Estimated Total Size (MB): 0.30				

FIGURE 5. Architecture of Lenet-5

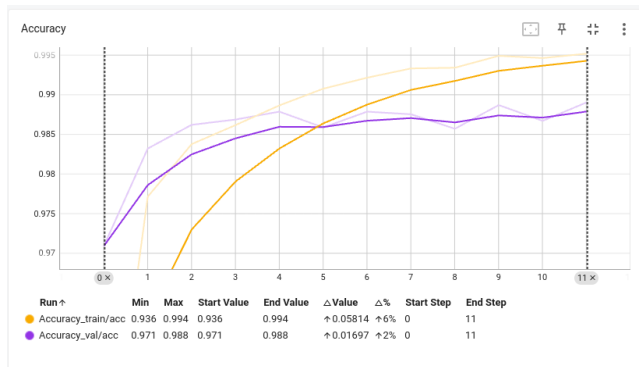


FIGURE 6. Accuracy Graph of Pytorch Lenet-5





FIGURE 7. Loss Graph of Pytorch Lenet-5






**FIGURE 8.** Loss Graph of Manual LeNet-5

```


TRAIN--> Correct: 54076 out of 60000, acc=0.9012666666666667
TEST--> Correct: 9027 out of 10000, acc=0.9027

```

**FIGURE 9.** Train Accuracy Vs Loss Accuracy of Manual LeNet-5



**FIGURE 10.** Predicted Output of LeNet-5



**FIGURE 11.** Predicted output Gradient View of LeNet-5

## B. APPENDIX B: CODE

```

1 import pickle
2 import random
3 import numpy as np
4 import matplotlib.pyplot as plt
5 from abc import ABCMeta, abstractmethod
6
7 filename = [
8     ["training_images", "train-images-idx3-ubyte.gz"],
9     ["test_images", "t10k-images-idx3-ubyte.gz"],
10    ["training_labels", "train-labels-idx1-ubyte.gz"],
11    ["test_labels", "t10k-labels-idx1-ubyte.gz"]
12 ]
13
14 def download_mnist():
15     base_url = "http://yann.lecun.com/exdb/mnist/"
16     for name in filename:
17         print("Downloading "+name[1]+"...")
18         request.urlretrieve(base_url+name[1], name[1])
19         print("Download complete.")
20
21 def save_mnist():
22     mnist = {}
23     for name in filename[2:]:
24         with gzip.open(name[1], 'rb') as f:
25             mnist[name[0]] = np.frombuffer(f.read(), np.uint8, offset=16).reshape(-1, 28*28)
26     for name in filename[-2:]:
27         with gzip.open(name[1], 'rb') as f:
28             mnist[name[0]] = np.frombuffer(f.read(), np.uint8, offset=8)
29     with open("mnist.pkl", 'wb') as f:
30         pickle.dump(mnist, f)
31     print("Save complete.")
32
33 def init():
34     download_mnist()
35     save_mnist()
36
37 def load():
38     with open("mnist.pkl", 'rb') as f:
39         mnist = pickle.load(f)
40     return mnist["training_images"], mnist["training_labels"], mnist["test_images"], mnist["test_labels"]
41
42 def MakeOneHot(Y, D_out):
43     N = Y.shape[0]
44     Z = np.zeros((N, D_out))
45     Z[np.arange(N), Y] = 1
46     return Z
47
48 def draw_losses(losses):
49     t = np.arange(len(losses))
50     plt.plot(t, losses)
51     plt.show()
52
53 def get_batch(X, Y, batch_size):
54     N = len(X)
55     i = random.randint(1, N-batch_size)
56     return X[i:i+batch_size], Y[i:i+batch_size]
57
58 class FC():
59     """
60     Fully connected layer
61     """
62     def __init__(self, D_in, D_out):
63         #print("Build FC")
64         self.cache = None
65         #self.W = {'val': np.random.randn(D_in,

```

```

D_out), 'grad': 0}
66         self.W = {'val': np.random.normal(0.0, np.
sqrt(2/D_in), (D_in, D_out)), 'grad': 0}
67         self.b = {'val': np.random.randn(D_out), '
grad': 0}
68
69     def _forward(self, X):
70         #print("FC: _forward")
71         out = np.dot(X, self.W['val']) + self.b['
val']
72         self.cache = X
73         return out
74
75     def _backward(self, dout):
76         #print("FC: _backward")
77         X = self.cache
78         dX = np.dot(dout, self.W['val'].T).reshape
(X.shape)
79         self.W['grad'] = np.dot(X.reshape(X.shape
[0], np.prod(X.shape[1:])).T, dout)
80         self.b['grad'] = np.sum(dout, axis=0)
81         #self._update_params()
82         return dX
83
84     def _update_params(self, lr=0.001):
85         # Update the parameters
86         self.W['val'] -= lr*self.W['grad']
87         self.b['val'] -= lr*self.b['grad']
88
89 class ReLU():
90     """
91     ReLU activation layer
92     """
93     def __init__(self):
94         #print("Build ReLU")
95         self.cache = None
96
97     def _forward(self, X):
98         #print("ReLU: _forward")
99         out = np.maximum(0, X)
100         self.cache = X
101         return out
102
103     def _backward(self, dout):
104         #print("ReLU: _backward")
105         X = self.cache
106         dX = np.array(dout, copy=True)
107         dX[X <= 0] = 0
108         return dX
109
110 class Sigmoid():
111     """
112     Sigmoid activation layer
113     """
114     def __init__(self):
115         self.cache = None
116
117     def _forward(self, X):
118         self.cache = X
119         return 1 / (1 + np.exp(-X))
120
121     def _backward(self, dout):
122         X = self.cache
123         dX = dout*X*(1-X)
124         return dX
125
126 class tanh():
127     """
128     tanh activation layer
129     """
130     def __init__(self):
131         self.cache = X
132
133     def _forward(self, X):

```

```

134     self.cache = X
135     return np.tanh(X)
136
137     def _backward(self, X):
138         X = self.cache
139         dX = dout*(1 - np.tanh(X)**2)
140         return dX
141
142 class Softmax():
143     """
144     Softmax activation layer
145     """
146     def __init__(self):
147         #print("Build Softmax")
148         self.cache = None
149
150     def _forward(self, X):
151         #print("Softmax: _forward")
152         maxes = np.amax(X, axis=1)
153         maxes = maxes.reshape(maxes.shape[0], 1)
154         Y = np.exp(X - maxes)
155         Z = Y / np.sum(Y, axis=1).reshape(Y.shape
156         [0], 1)
157         self.cache = (X, Y, Z)
158         return Z # distribution
159
160     def _backward(self, dout):
161         X, Y, Z = self.cache
162         dZ = np.zeros(X.shape)
163         dY = np.zeros(X.shape)
164         dX = np.zeros(X.shape)
165         N = X.shape[0]
166         for n in range(N):
167             i = np.argmax(Z[n])
168             dZ[n,:] = np.diag(Z[n]) - np.outer(Z[n
169             ],Z[n])
170             M = np.zeros((N,N))
171             M[:,i] = 1
172             dY[n,:] = np.eye(N) - M
173             dX = np.dot(dout,dZ)
174             dX = np.dot(dX,dY)
175             return dX
176
177 class Dropout():
178     """
179     Dropout layer
180     """
181     def __init__(self, p=1):
182         self.cache = None
183         self.p = p
184
185     def _forward(self, X):
186         M = (np.random.rand(*X.shape) < self.p) /
187         self.p
188         self.cache = X, M
189         return X*M
190
191     def _backward(self, dout):
192         X, M = self.cache
193         dX = dout*M/self.p
194         return dX
195
196 class Conv():
197     """
198     Conv layer
199     """
200     def __init__(self, Cin, Cout, F, stride=1,
201     padding=0, bias=True):
202         self.Cin = Cin
203         self.Cout = Cout
204         self.F = F
205         self.S = stride
206         #self.W = {'val': np.random.randn(Cout,
207         Cin, F, F), 'grad': 0}
208
209         self.W = {'val': np.random.normal(0.0,np.
210         sqrt(2/Cin), (Cout,Cin,F,F)), 'grad': 0} #
211         Xavier Initialization
212         self.b = {'val': np.random.randn(Cout), '
213         grad': 0}
214         self.cache = None
215         self.pad = padding
216
217     def _forward(self, X):
218         X = np.pad(X, ((0,0), (0,0), (self.pad,self.
219         pad)), (self.pad,self.pad)), 'constant')
220         (N, Cin, H, W) = X.shape
221         H_ = H - self.F + 1
222         W_ = W - self.F + 1
223         Y = np.zeros((N, self.Cout, H_, W_))
224
225         for n in range(N):
226             for c in range(self.Cout):
227                 for h in range(H_):
228                     for w in range(W_):
229                         Y[n, c, h, w] = np.sum(X[n
230                         , :, h:h+self.F, w:w+self.F] * self.W['val'])[c
231                         , :, :, :]) + self.b['val'][c]
232
233         self.cache = X
234         return Y
235
236     def _backward(self, dout):
237         # dout (N,Cout,H_,W_)
238         # W (Cout, Cin, F, F)
239         X = self.cache
240         (N, Cin, H, W) = X.shape
241         H_ = H - self.F + 1
242         W_ = W - self.F + 1
243         W_rot = np.rot90(np.rot90(self.W['val']))
244
245         dX = np.zeros(X.shape)
246         dW = np.zeros(self.W['val'].shape)
247         db = np.zeros(self.b['val'].shape)
248
249         # dW
250         for co in range(self.Cout):
251             for ci in range(Cin):
252                 for h in range(self.F):
253                     for w in range(self.F):
254                         dW[co, ci, h, w] = np.sum(
255                         X[:,ci,h:h+W_,w:w+W_] * dout[:,co,:,:])
256
257         # db
258         for co in range(self.Cout):
259             db[co] = np.sum(dout[:,co,:,:])
260
261         dout_pad = np.pad(dout, ((0,0), (0,0), (self
262         .F,self.F), (self.F,self.F)), 'constant')
263         #print("dout_pad.shape: " + str(dout_pad.
264         shape))
265         # dX
266         for n in range(N):
267             for ci in range(Cin):
268                 for h in range(H):
269                     for w in range(W):
270                         #print("self.F.shape: %s",
271                         self.F)
272                         #print("%s, W_rot[:,ci
273                         ,:,:].shape: %s, dout_pad[n,:,h:h+self.F,w:w+
274                         self.F].shape: %s" % ((n,ci,h,w),W_rot[:,ci
275                         ,:,:].shape, dout_pad[n,:,h:h+self.F,w:w+self.
276                         F].shape))
277                         dX[n, ci, h, w] = np.sum(
278                         W_rot[:,ci,:,:) * dout_pad[n, :, h:h+self.F,w:
279                         w+self.F])
280
281         return dX

```

```

261 class MaxPool():
262     def __init__(self, F, stride):
263         self.F = F
264         self.S = stride
265         self.cache = None
266
267     def _forward(self, X):
268         # X: (N, Cin, H, W): maxpool along 3rd, 4
269         th dim
270         (N,Cin,H,W) = X.shape
271         F = self.F
272         W_ = int(float(W)/F)
273         H_ = int(float(H)/F)
274         Y = np.zeros((N,Cin,W_,H_))
275         M = np.zeros(X.shape) # mask
276         for n in range(N):
277             for cin in range(Cin):
278                 for w_ in range(W_):
279                     for h_ in range(H_):
280                         Y[n,cin,w_,h_] = np.max(X[
281                             n,cin,F*w_:F*(w_+1),F*h_:F*(h_+1)])
282                         i,j = np.unravel_index(X[n
283                             ,cin,F*w_:F*(w_+1),F*h_:F*(h_+1)].argmax(), (F
284                             ,F))
285                         M[n,cin,F*w_+i,F*h_+j] = 1
286         self.cache = M
287         return Y
288
289     def _backward(self, dout):
290         M = self.cache
291         (N,Cin,H,W) = M.shape
292         dout = np.array(dout)
293         #print("dout.shape: %s, M.shape: %s" % (
294             dout.shape, M.shape))
295         dX = np.zeros(M.shape)
296         for n in range(N):
297             for c in range(Cin):
298                 #print("(n,c): (%s,%s)" % (n,c))
299                 dX[n,c,:,:] = dout[n,c,:,:].repeat
300                 (2, axis=0).repeat(2, axis=1)
301             return dX*M
302
303 def NLLLoss(Y_pred, Y_true):
304     """
305     Negative log likelihood loss
306     """
307     loss = 0.0
308     N = Y_pred.shape[0]
309     M = np.sum(Y_pred*Y_true, axis=1)
310     for e in M:
311         #print(e)
312         if e == 0:
313             loss += 500
314         else:
315             loss += -np.log(e)
316     return loss/N
317
318 class CrossEntropyLoss():
319     def __init__(self):
320         pass
321
322     def get(self, Y_pred, Y_true):
323         N = Y_pred.shape[0]
324         softmax = Softmax()
325         prob = softmax._forward(Y_pred)
326         loss = NLLLoss(prob, Y_true)
327         Y_serial = np.argmax(Y_true, axis=1)
328         dout = prob.copy()
329         dout[np.arange(N), Y_serial] -= 1
330         return loss, dout
331
332 class SoftmaxLoss():
333     def __init__(self):
334         pass

```

```

329
330     def get(self, Y_pred, Y_true):
331         N = Y_pred.shape[0]
332         loss = NLLLoss(Y_pred, Y_true)
333         Y_serial = np.argmax(Y_true, axis=1)
334         dout = Y_pred.copy()
335         dout[np.arange(N), Y_serial] -= 1
336         return loss, dout
337
338 class Net(metaclass=ABCMeta):
339     # Neural network super class
340
341     @abstractmethod
342     def __init__(self):
343         pass
344
345     @abstractmethod
346     def forward(self, X):
347         pass
348
349     @abstractmethod
350     def backward(self, dout):
351         pass
352
353     @abstractmethod
354     def get_params(self):
355         pass
356
357     @abstractmethod
358     def set_params(self, params):
359         pass
360
361 class TwoLayerNet(Net):
362     #Simple 2 layer NN
363
364     def __init__(self, N, D_in, H, D_out, weights=
365         ''):
366         self.FC1 = FC(D_in, H)
367         self.ReLU1 = ReLU()
368         self.FC2 = FC(H, D_out)
369
370         if weights == '':
371             pass
372         else:
373             with open(weights,'rb') as f:
374                 params = pickle.load(f)
375                 self.set_params(params)
376
377     def forward(self, X):
378         h1 = self.FC1._forward(X)
379         a1 = self.ReLU1._forward(h1)
380         h2 = self.FC2._forward(a1)
381         return h2
382
383     def backward(self, dout):
384         dout = self.FC2._backward(dout)
385         dout = self.ReLU1._backward(dout)
386         dout = self.FC1._backward(dout)
387
388     def get_params(self):
389         return [self.FC1.W, self.FC1.b, self.FC2.W
390             , self.FC2.b]
391
392     def set_params(self, params):
393         [self.FC1.W, self.FC1.b, self.FC2.W, self.
394             FC2.b] = params
395
396 class ThreeLayerNet(Net):
397     #Simple 3 layer NN

```

```

400 def __init__(self, N, D_in, H1, H2, D_out,
401 weights=''):
402     self.FC1 = FC(D_in, H1)
403     self.ReLU1 = ReLU()
404     self.FC2 = FC(H1, H2)
405     self.ReLU2 = ReLU()
406     self.FC3 = FC(H2, D_out)
407
408     if weights == '':
409         pass
410     else:
411         with open(weights, 'rb') as f:
412             params = pickle.load(f)
413             self.set_params(params)
414
415 def forward(self, X):
416     h1 = self.FC1._forward(X)
417     a1 = self.ReLU1._forward(h1)
418     h2 = self.FC2._forward(a1)
419     a2 = self.ReLU2._forward(h2)
420     h3 = self.FC3._forward(a2)
421     return h3
422
423 def backward(self, dout):
424     dout = self.FC3._backward(dout)
425     dout = self.ReLU2._backward(dout)
426     dout = self.FC2._backward(dout)
427     dout = self.ReLU1._backward(dout)
428     dout = self.FC1._backward(dout)
429
430 def get_params(self):
431     return [self.FC1.W, self.FC1.b, self.FC2.W,
432             self.FC2.b, self.FC3.W, self.FC3.b]
433
434 def set_params(self, params):
435     [self.FC1.W, self.FC1.b, self.FC2.W, self.
436      FC2.b, self.FC3.W, self.FC3.b] = params
437
438 class LeNet5(Net):
439     # LeNet5
440
441     def __init__(self):
442         self.conv1 = Conv(1, 6, 5)
443         self.ReLU1 = ReLU()
444         self.pool1 = MaxPool(2,2)
445         self.conv2 = Conv(6, 16, 5)
446         self.ReLU2 = ReLU()
447         self.pool2 = MaxPool(2,2)
448         self.FC1 = FC(16*4*4, 120)
449         self.ReLU3 = ReLU()
450         self.FC2 = FC(120, 84)
451         self.ReLU4 = ReLU()
452         self.FC3 = FC(84, 10)
453         self.Softmax = Softmax()
454
455         self.p2_shape = None
456
457     def forward(self, X):
458         h1 = self.conv1._forward(X)
459         a1 = self.ReLU1._forward(h1)
460         p1 = self.pool1._forward(a1)
461         h2 = self.conv2._forward(p1)
462         a2 = self.ReLU2._forward(h2)
463         p2 = self.pool2._forward(a2)
464         self.p2_shape = p2.shape
465         f1 = p2.reshape(X.shape[0], -1) # Flatten
466         h3 = self.FC1._forward(f1)
467         a3 = self.ReLU3._forward(h3)
468         h4 = self.FC2._forward(a3)
469         a5 = self.ReLU4._forward(h4)
470         h5 = self.FC3._forward(a5)
471         a5 = self.Softmax._forward(h5)
472         return a5
473
474 def backward(self, dout):
475     #dout = self.Softmax._backward(dout)
476     dout = self.FC3._backward(dout)
477     dout = self.ReLU4._backward(dout)
478     dout = self.FC2._backward(dout)
479     dout = self.ReLU3._backward(dout)
480     dout = self.FC1._backward(dout)
481     dout = dout.reshape(self.p2_shape) #
482     reshape
483     dout = self.pool2._backward(dout)
484     dout = self.ReLU2._backward(dout)
485     dout = self.conv2._backward(dout)
486     dout = self.pool1._backward(dout)
487     dout = self.ReLU1._backward(dout)
488     dout = self.conv1._backward(dout)
489
490 def get_params(self):
491     return [self.conv1.W, self.conv1.b, self.
492             conv2.W, self.conv2.b, self.FC1.W, self.FC1.b,
493             self.FC2.W, self.FC2.b, self.FC3.W, self.FC3.
494             b]
495
496 def set_params(self, params):
497     [self.conv1.W, self.conv1.b, self.conv2.W,
498      self.conv2.b, self.FC1.W, self.FC1.b, self.
499      FC2.W, self.FC2.b, self.FC3.W, self.FC3.b] =
500     params
501
502 class SGD():
503     def __init__(self, params, lr=0.001, reg=0):
504         self.parameters = params
505         self.lr = lr
506         self.reg = reg
507
508     def step(self):
509         for param in self.parameters:
510             param['val'] -= (self.lr*param['grad']
511                             + self.reg*param['val'])
512
513 class SGDMomentum():
514     def __init__(self, params, lr=0.001, momentum
515                  =0.99, reg=0):
516         self.l = len(params)
517         self.parameters = params
518         self.velocities = []
519         for param in self.parameters:
520             self.velocities.append(np.zeros(param[
521             'val'].shape))
522         self.lr = lr
523         self.rho = momentum
524         self.reg = reg
525
526     def step(self):
527         for i in range(self.l):
528             self.velocities[i] = self.rho*self.
529             velocities[i] + (1-self.rho)*self.parameters[i
530             ]['grad']
531             self.parameters[i]['val'] -= (self.lr*
532             self.velocities[i] + self.reg*self.parameters[i
533             ]['val'])
534
535 """
536 (1) Prepare Data: Load, Shuffle, Normalization,
537 Batching, Preprocessing
538 """
539
540 #mnist.init()
541 X_train, Y_train, X_test, Y_test = load()
542 X_train, X_test = X_train/float(255), X_test/float
543 (255)
544 X_train -= np.mean(X_train)
545 X_test -= np.mean(X_test)

```



```

529
530 batch_size = 64
531 D_in = 784
532 D_out = 10
533
534 print("batch_size: " + str(batch_size) + ", D_in: "
      + str(D_in) + ", D_out: " + str(D_out))
535
536 ### TWO LAYER NET FORWARD TEST ###
537 #H=400
538 #model = nn.TwoLayerNet(batch_size, D_in, H, D_out
539 )
540 H1=300
541 H2=100
542 model = ThreeLayerNet(batch_size, D_in, H1, H2,
543 D_out)
544
545 losses = []
546 #optim = optimizer.SGD(model.get_params(), lr
547 =0.0001, reg=0)
548 optim = SGDMomentum(model.get_params(), lr=0.0001,
549 momentum=0.80, reg=0.00003)
550 criterion = CrossEntropyLoss()
551
552 # TRAIN
553 ITER = 2500
554 for i in range(ITER):
555     # get batch, make onehot
556     X_batch, Y_batch = get_batch(X_train, Y_train,
557 batch_size)
558     Y_batch = MakeOneHot(Y_batch, D_out)
559
560     # forward, loss, backward, step
561     Y_pred = model.forward(X_batch)
562     loss, dout = criterion.get(Y_pred, Y_batch)
563     model.backward(dout)
564     optim.step()
565
566     if i % 10 == 0:
567         print("%s%% iter: %s, loss: %s" % (10*i/ITER, i
568 , loss))
569         losses.append(loss)
570
571 # save params
572 weights = model.get_params()
573 with open("weights.pkl", "wb") as f:
574     pickle.dump(weights, f)
575
576 draw_losses(losses)
577
578 # TRAIN SET ACC
579 Y_pred = model.forward(X_train)
580 result = np.argmax(Y_pred, axis=1) - Y_train
581 result = list(result)
582 print("TRAIN--> Correct: " + str(result.count(0))
583 + " out of " + str(X_train.shape[0]) + ", acc="
584 + str(result.count(0)/X_train.shape[0]))
585
586 # TEST SET ACC
587 Y_pred = model.forward(X_test)
588 result = np.argmax(Y_pred, axis=1) - Y_test
589 result = list(result)
590 print("TEST--> Correct: " + str(result.count(0)) +
591 " out of " + str(X_test.shape[0]) + ", acc="
592 + str(result.count(0)/X_test.shape[0]))

```

...